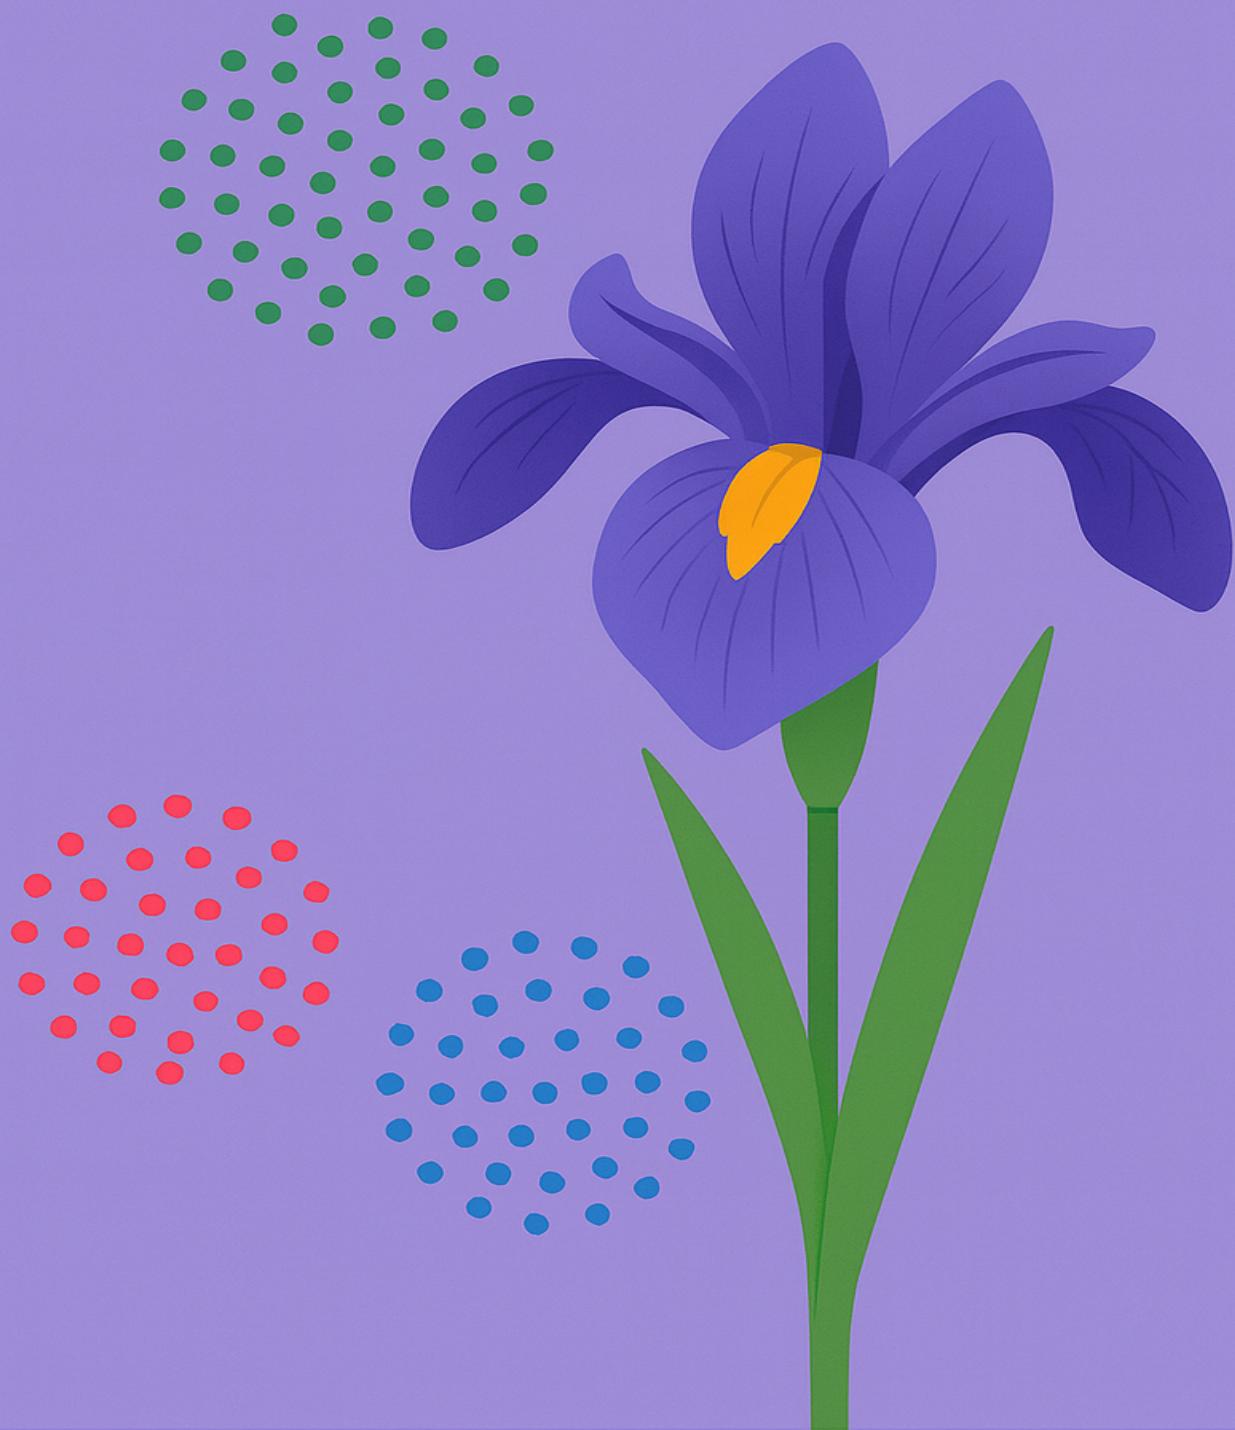


Iris Clusters Analysis



Iris Clusters Analysis

Seif H. Kungulio

December 03, 2025

Contents

Business Understanding	1
Introduction	1
Problem Statement	1
Data Understanding	2
Load the Dataset	2
Statistical Summary	2
Data Description	3
Data Dictionary	3
Check for Missing Values	3
Exploratory Data Analysis (Plots)	3
Pairwise relationship	3
Univariate distribution	4
Data Preparation	6
Select Features and Standardize	6
PCA for Dimensionality Reduction & Visualization	6
Modeling	8
Determine number of clusters (k) for K-Means	8
K-Means Clustering	9
Hierarchical Clustering	11
Gaussian Mixture Model (Model-Based Clustering)	12
Evaluation	14
Confusion Tables	14
K-Means vs Species	14
Hierarchical vs Species	14
GMM vs Species	14
Adjusted Rand Index (ARI)	14
Silhouette Analysis for K-Means	15
Deployment	17
Conclusion	18
About the Author	19

Business Understanding

Introduction

Problem Statement

Data Understanding

Load the Dataset

Load the dataset and display first six rows of the dataset

```
data(iris)
head(iris)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1          5.1         3.5          1.4         0.2  setosa
## 2          4.9         3.0          1.4         0.2  setosa
## 3          4.7         3.2          1.3         0.2  setosa
## 4          4.6         3.1          1.5         0.2  setosa
## 5          5.0         3.6          1.4         0.2  setosa
## 6          5.4         3.9          1.7         0.4  setosa
```

Statistical Summary

Display the statistical summary of the dataset

```
summary(iris)
```

```
##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
## Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100
## 1st Qu.:5.100  1st Qu.:2.800  1st Qu.:1.600  1st Qu.:0.300
## Median :5.800  Median :3.000  Median :4.350  Median :1.300
## Mean   :5.843  Mean   :3.057  Mean   :1.758  Mean   :1.199
## 3rd Qu.:6.400  3rd Qu.:3.300  3rd Qu.:5.100  3rd Qu.:1.800
## Max.   :7.900  Max.   :4.400  Max.   :6.900  Max.   :2.500
## 
##   Species
##   setosa      :50
##   versicolor :50
##   virginica  :50
## 
## 
##
```

```
skim(iris)
```

Table 1: Data summary

Name	iris
Number of rows	150
Number of columns	5
Column type frequency:	
factor	1
numeric	4

Group variables	None
-----------------	------

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
Species	0	1	FALSE	3	set: 50, ver: 50, vir: 50

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
Sepal.Length	0	1	5.84	0.83	4.3	5.1	5.80	6.4	7.9	
Sepal.Width	0	1	3.06	0.44	2.0	2.8	3.00	3.3	4.4	
Petal.Length	0	1	3.76	1.77	1.0	1.6	4.35	5.1	6.9	
Petal.Width	0	1	1.20	0.76	0.1	0.3	1.30	1.8	2.5	

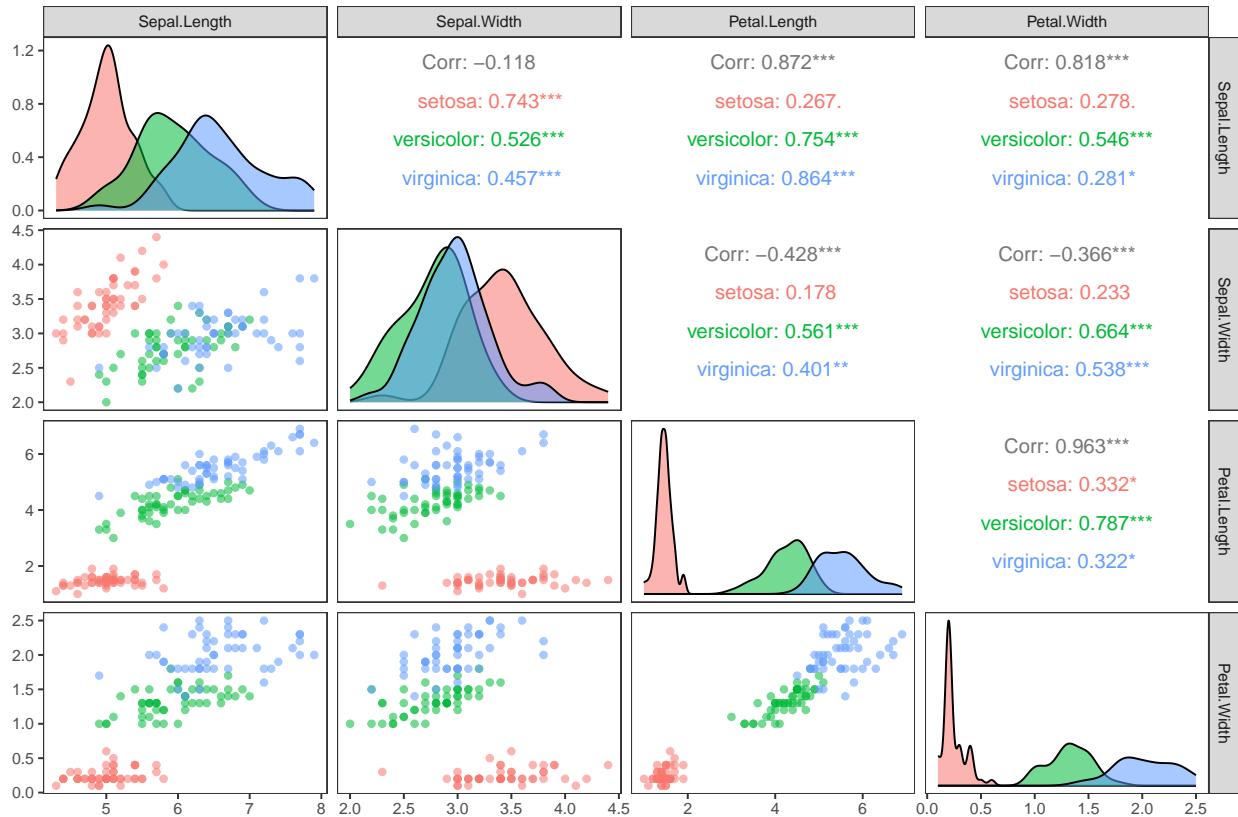
Data Description**Data Dictionary****Check for Missing Values**

```
colSums(is.na(iris))
```

```
## Sepal.Length  Sepal.Width  Petal.Length  Petal.Width      Species
##          0          0          0          0          0
```

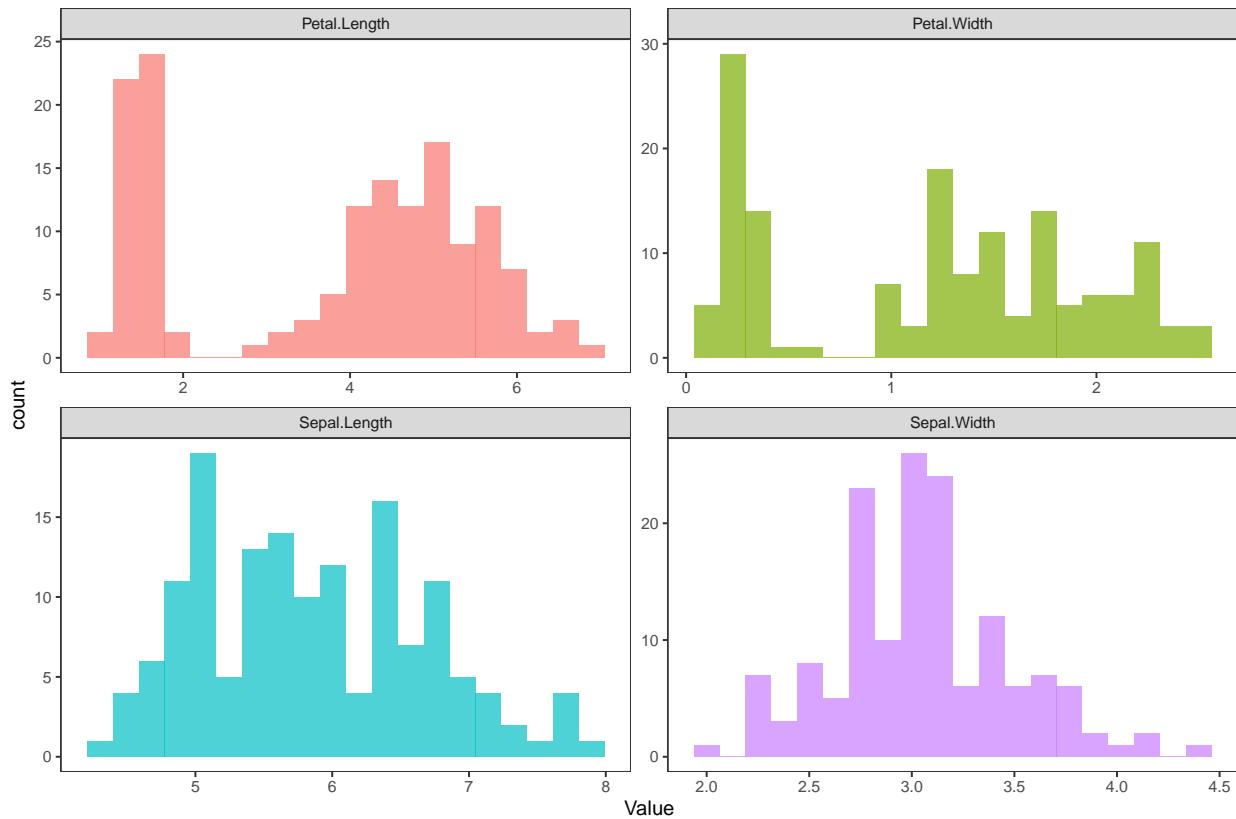
Exploratory Data Analysis (Plots)**Pairwise relationship**

```
ggpairs(
  iris,
  columns = 1:4,
  aes(color = Species,
      alpha = 0.7
    )
)
```



Univariate distribution

```
iris %>%
  pivot_longer(cols = 1:4, names_to = "Feature", values_to = "Value") %>%
  ggplot(aes(x = Value, fill = Feature)) +
  geom_histogram(bins = 20, alpha = 0.7, show.legend = FALSE) +
  facet_wrap(~ Feature, scale = "free")
```



Data Preparation

Select Features and Standardize

For unsupervised modeling, I'll use only the numeric features and standardize them, so that variables with larger scales don't dominate distance calculations.

Select numeric features only

```
iris_features <- iris %>%
  select(Sepal.Length, Sepal.Width, Petal.Length, Petal.Width)
```

Standardize/ scale the numeric features

```
iris_scaled <- scale(iris_features)
```

Quick check of the standardized numeric features

```
summary(iris_scaled)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## Min.	-1.86378	-2.4258	-1.5623	-1.4422
## 1st Qu.	-0.89767	-0.5904	-1.2225	-1.1799
## Median	-0.05233	-0.1315	0.3354	0.1321
## Mean	0.00000	0.0000	0.0000	0.0000
## 3rd Qu.	0.67225	0.5567	0.7602	0.7880
## Max.	2.48370	3.0805	1.7799	1.7064

PCA for Dimensionality Reduction & Visualization

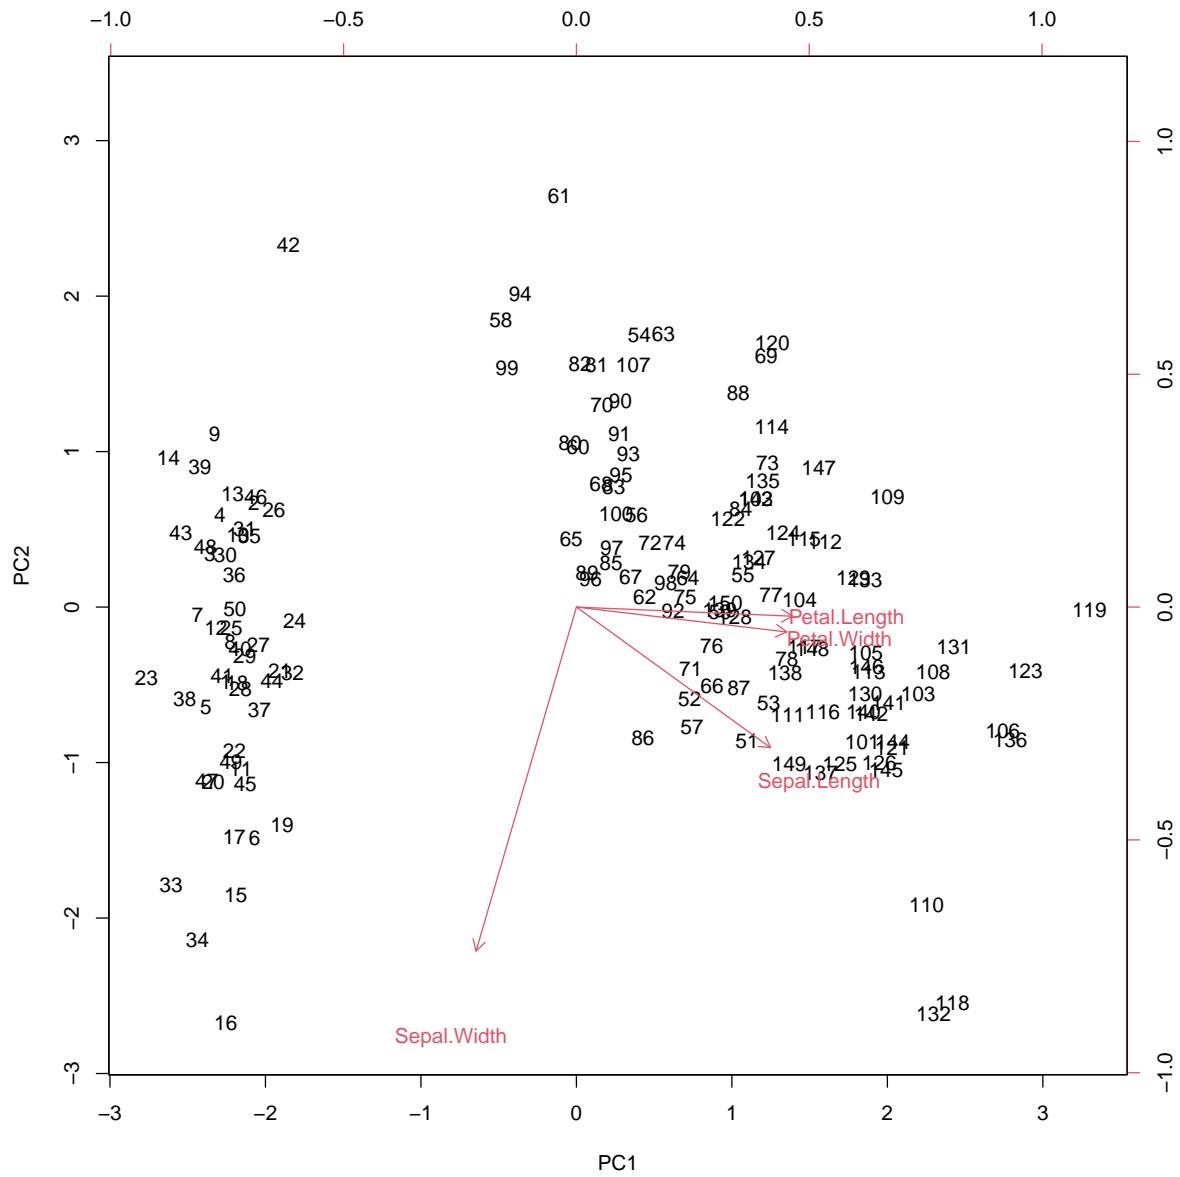
```
pca_model <- prcomp(iris_scaled, center = TRUE, scale = TRUE)

summary(pca_model)
```

	PC1	PC2	PC3	PC4
## Importance of components:				
## Standard deviation	1.7084	0.9560	0.38309	0.14393
## Proportion of Variance	0.7296	0.2285	0.03669	0.00518
## Cumulative Proportion	0.7296	0.9581	0.99482	1.00000

Biplot

```
biplot(pca_model, scale = 0)
```



Modeling

I will explore several unsupervised methods:

- K-means clustering
- Hierarchical clustering
- Gaussian mixture models (GMM) via mclust

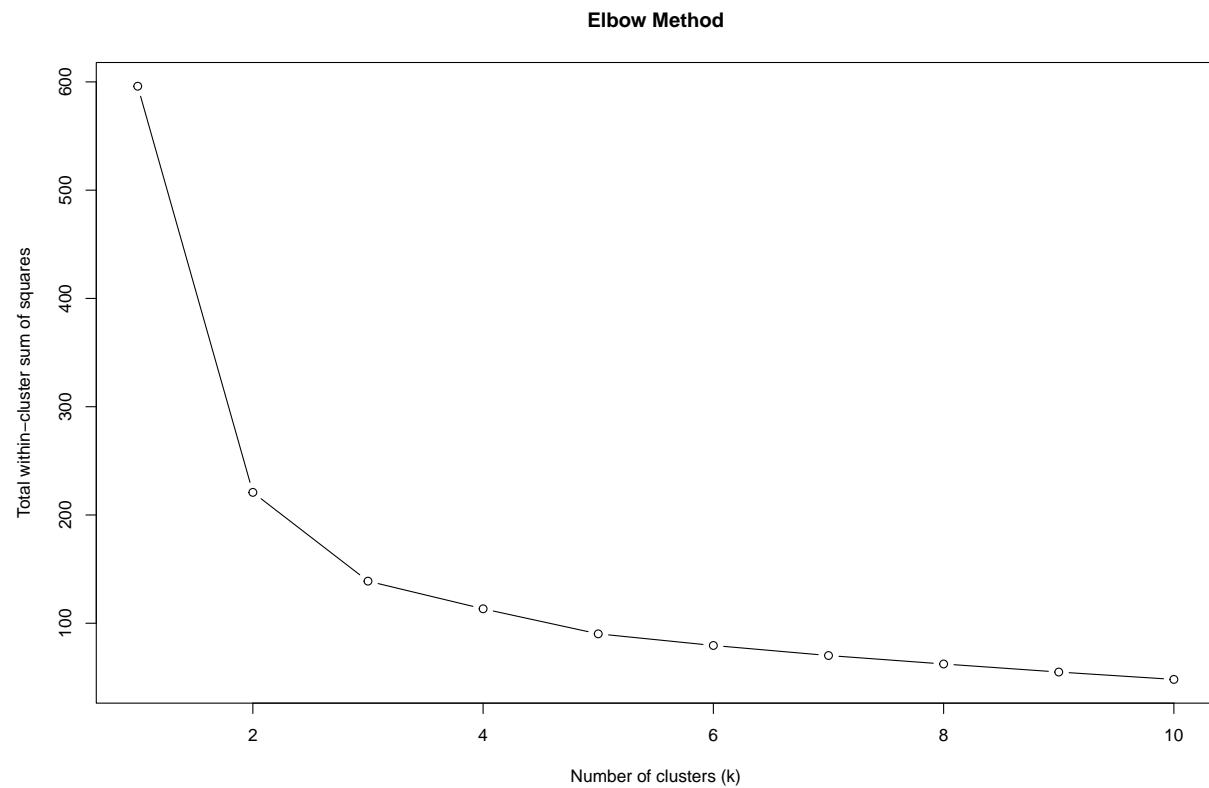
Determine number of clusters (k) for K-Means

Elbow Method

```
set.seed(123)

wss <- map_dbl(1:10, ~ {
  kmeans(iris_scaled, centers = .x, nstart = 25)$tot.withinss
})

plot(
  1:10, wss, type = "b",
  xlab = "Number of clusters (k)",
  ylab = "Total within-cluster sum of squares",
  main = "Elbow Method"
)
```



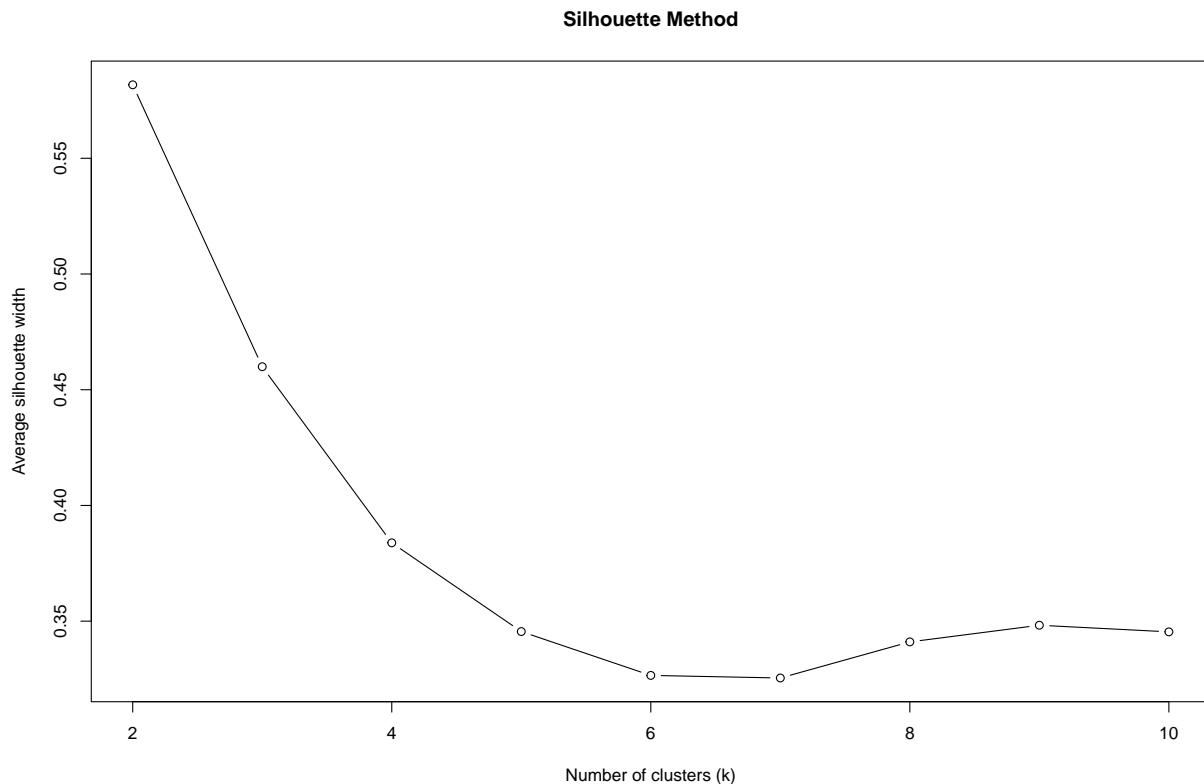
Average Silhouette Width

```

sil_width <- map_dbl(2:10, ~ {
  km <- kmeans(iris_scaled, centers = .x, nstart = 25)
  ss <- silhouette(km$cluster, dist(iris_scaled))
  mean(ss[, "sil_width"])
})

plot(
  2:10, sil_width, type = "b",
  xlab = "Number of clusters (k)",
  ylab = "Average silhouette width",
  main = "Silhouette Method"
)

```



K-Means Clustering

A partitioning method that divides the dataset into k clusters based on distance to centroids.

```

set.seed(123)

k_opt <- 3 # chosen after inspecting elbow/silhouette
kmeans_model <- kmeans(iris_scaled, centers = k_opt, nstart = 25)

kmeans_model$size

## [1] 50 53 47

```

```
kmeans_model$centers
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width
## 1 -1.01119138  0.85041372 -1.3006301 -1.2507035
## 2 -0.05005221 -0.88042696  0.3465767  0.2805873
## 3  1.13217737  0.08812645  0.9928284  1.0141287
```

Add cluster to original data

```
iris_kmeans <- iris %>%
  mutate(KMeansCluster = factor(kmeans_model$cluster))
head(iris_kmeans)
```

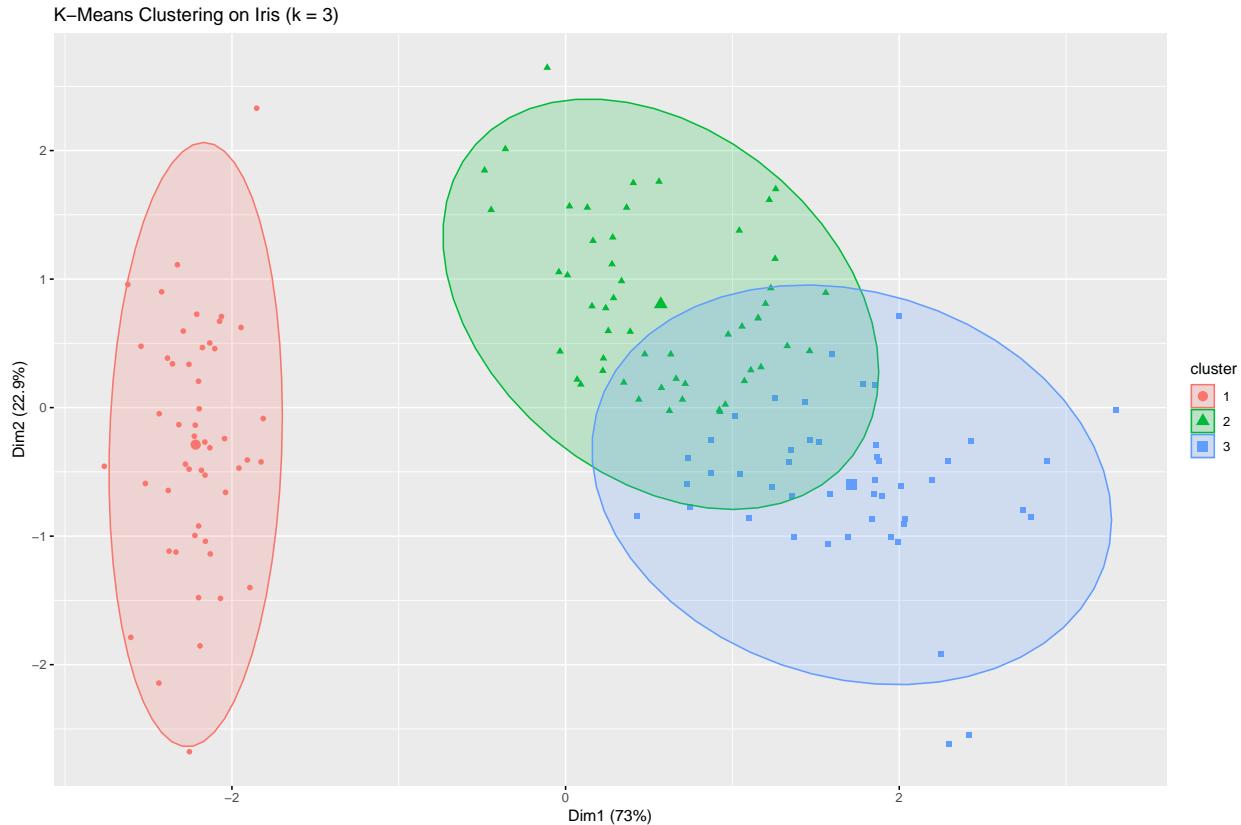
```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species KMeansCluster
## 1          5.1         3.5         1.4        0.2  setosa           1
## 2          4.9         3.0         1.4        0.2  setosa           1
## 3          4.7         3.2         1.3        0.2  setosa           1
## 4          4.6         3.1         1.5        0.2  setosa           1
## 5          5.0         3.6         1.4        0.2  setosa           1
## 6          5.4         3.9         1.7        0.4  setosa           1
```

```
tail(iris_kmeans)
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species KMeansCluster
## 145          6.7         3.3         5.7        2.5 virginica           3
## 146          6.7         3.0         5.2        2.3 virginica           3
## 147          6.3         2.5         5.0        1.9 virginica           2
## 148          6.5         3.0         5.2        2.0 virginica           3
## 149          6.2         3.4         5.4        2.3 virginica           3
## 150          5.9         3.0         5.1        1.8 virginica           2
```

Visualize K-Means clusters

```
fviz_cluster(
  kmeans_model,
  data = iris_scaled,
  geom = "point",
  ellipse.type = "norm",
  main = "K-Means Clustering on Iris (k = 3)"
)
```



Hierarchical Clustering

Builds a dendrogram to represent nested groupings

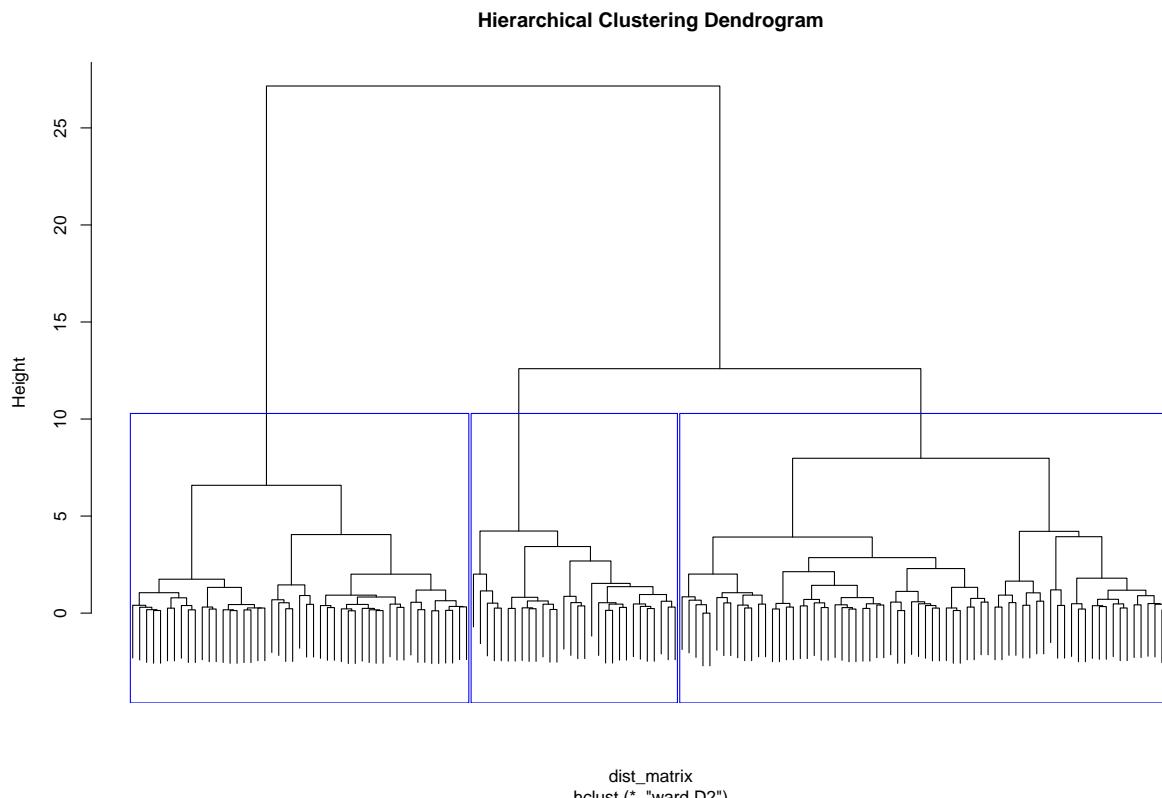
Create distance matrix

```
dist_matrix <- dist(iris_scaled)
```

Hierarchical clustering with Ward's method

```
hc_model <- hclust(dist_matrix, method = "ward.D2")

plot(hc_model, labels = FALSE, main = "Hierarchical Clustering Dendrogram")
rect.hclust(hc_model, k = 3, border = "blue")
```



Cut tree into 3 clusters

```
hc_clusters <- cutree(hc_model, k = 3)
iris_hclust <- iris %>%
  mutate(HCluster = factor(hc_clusters))

table(iris_hclust$HCluster)
```

```
##  
## 1 2 3  
## 49 30 71
```

Gaussian Mixture Model (Model-Based Clustering)

Assumes data comes from a mixture of Gaussian distributions.

Gaussian mixture models (GMM) via mclust

```
set.seed(123)

gmm_model <- Mclust(iris_scaled)

summary(gmm_model)

## -----
## Gaussian finite mixture model fitted by EM algorithm
```

```
## -----
## 
## Mclust VVV (ellipsoidal, varying volume, shape, and orientation) model with 2
## components:
## 
##   log-likelihood   n df      BIC      ICL
##   -322.6936 150 29 -790.6956 -790.6969
## 
## Clustering table:
##   1   2
## 50 100
```

Cluster assignment

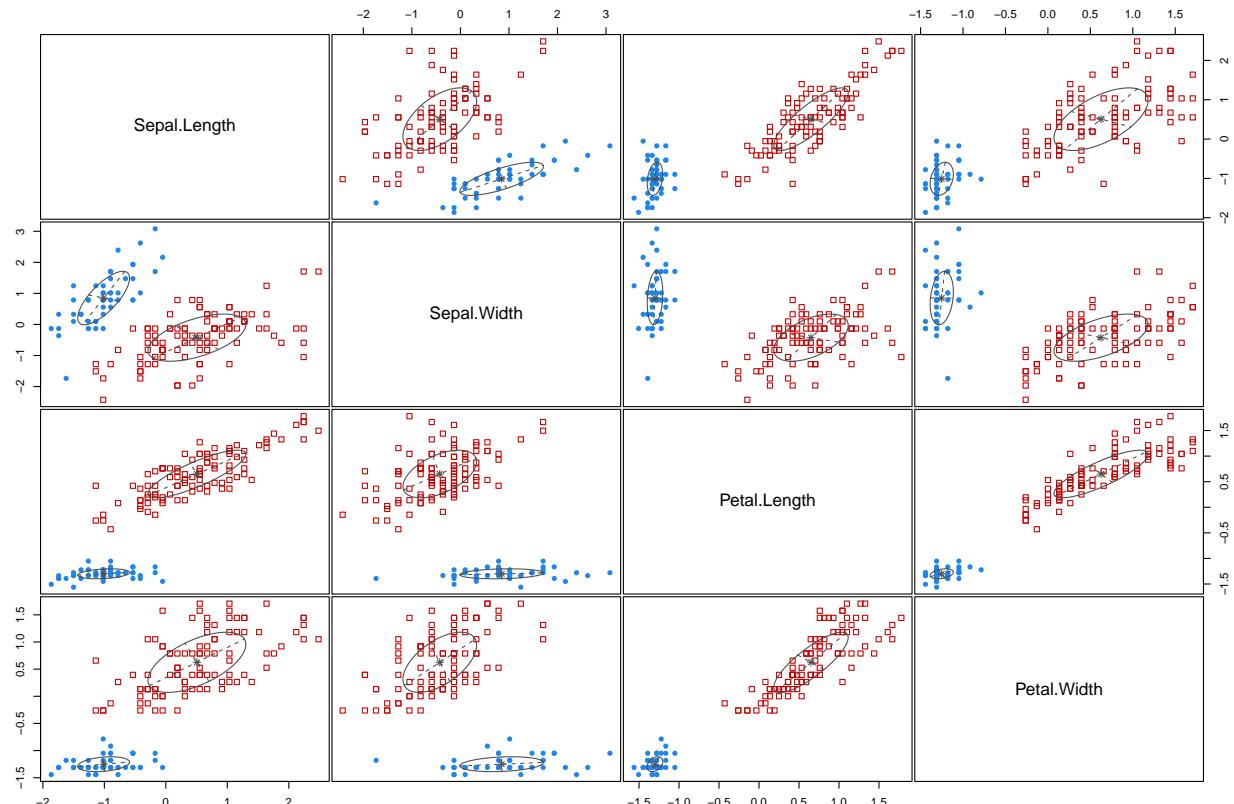
```
gmm_clusters <- gmm_model$classification
iris_gmm <- iris %>%
  mutate(GMMCcluster = factor(gmm_clusters))

table(iris_gmm$GMMCcluster)
```

```
## 
##   1   2
## 50 100
```

Plot GMM Clusters

```
plot(gmm_model, what = "classification")
```



Evaluation

Although these are unsupervised models, I can compare the resulting clusters with the known Species for evaluation purposes only.

Confusion Tables

K-Means vs Species

```
table(Cluster = iris_kmeans$KMeansCluster, Species = iris_kmeans$Species)
```

```
##          Species
## Cluster setosa versicolor virginica
##      1     50          0         0
##      2     0          39        14
##      3     0          11        36
```

Hierarchical vs Species

```
table(Cluster = iris_hclust$HCluster, Species = iris_hclust$Species)
```

```
##          Species
## Cluster setosa versicolor virginica
##      1     49          0         0
##      2     1          27        2
##      3     0          23        48
```

GMM vs Species

```
table(Cluster = iris_gmm$GMMCluster, Species = iris_gmm$Species)
```

```
##          Species
## Cluster setosa versicolor virginica
##      1     50          0         0
##      2     0          50        50
```

Adjusted Rand Index (ARI)

The Adjusted Rand Index measures agreement between two partitions (here: clusters vs. species), adjusted for chance.

Helper function to compute ARI

```
compute_ari <- function(cluster_labels) {
  cluster.stats(
    d = dist(iris_scaled),
    clustering = cluster_labels,
    alt.clustering = as.numeric(iris$Species)
  )$corrected.rand
}
```

Compute ARI

```
ari_kmeans <- compute_ari(kmeans_model$cluster)
ari_hclust <- compute_ari(hc_clusters)
ari_gmm <- compute_ari(gmm_clusters)
```

Display models ARI

```
ari_kmeans
```

```
## [1] 0.6201352
```

```
ari_hclust
```

```
## [1] 0.615323
```

```
ari_gmm
```

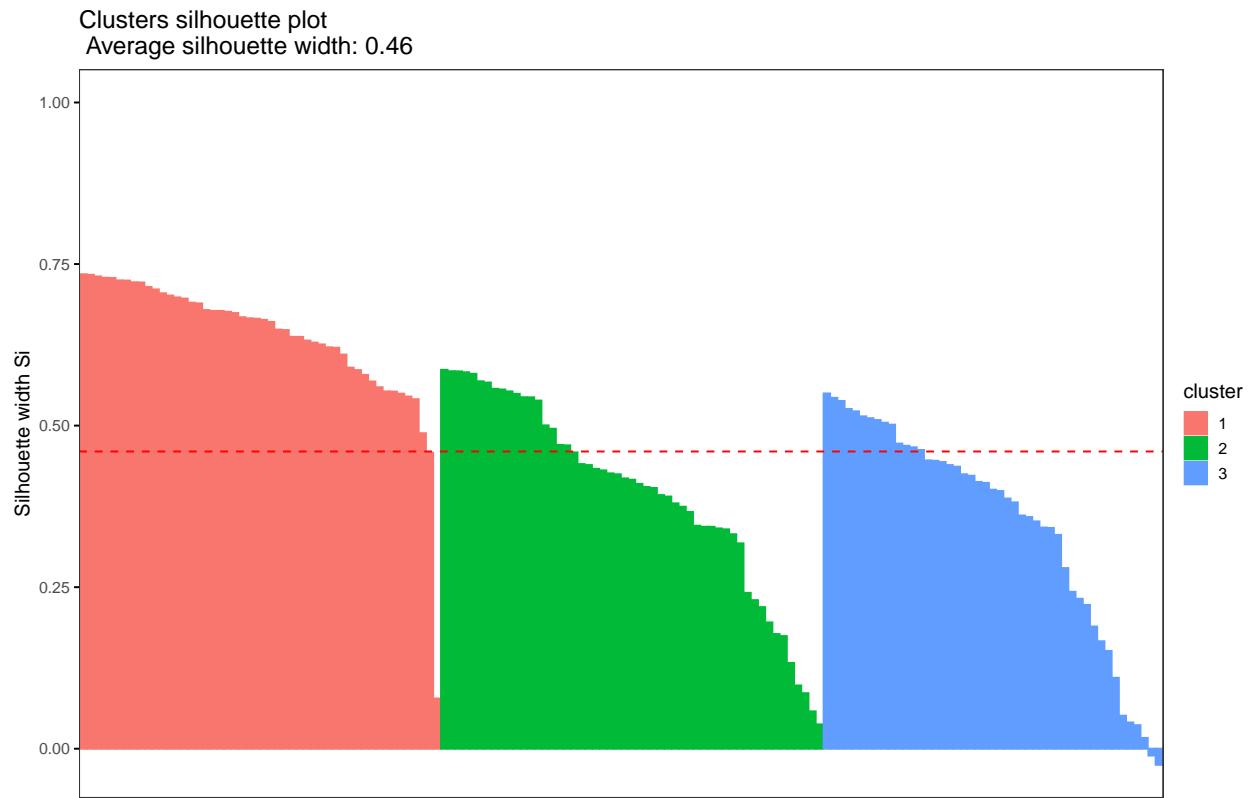
```
## [1] 0.5681159
```

Silhouette Analysis for K-Means

```
sil_kmeans <- silhouette(kmeans_model$cluster, dist(iris_scaled))
fviz_silhouette(sil_kmeans)
```

```
## Warning: `aes_string()` was deprecated in ggplot2 3.0.0.
## i Please use tidy evaluation idioms with `aes()``.
## i See also `vignette("ggplot2-in-packages")` for more information.
## i The deprecated feature was likely used in the factoextra package.
## Please report the issue at <https://github.com/kassambara/factoextra/issues>.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.

##   cluster size ave.sil.width
## 1       1   50      0.64
## 2       2   53      0.39
## 3       3   47      0.35
```



Deployment

Conclusion

About the Author